# Optimization for Machine Learning CS-439

Lecture 11: Gradient free and adaptive methods

**Nicolas Flammarion** 

EPFL - github.com/epfml/0ptML\_course May 16, 2025

# Chapter XI.1

# **Zero-Order Optimization**

#### Look mom no gradients!

Can we optimize  $\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x})$  without access to gradients?

meet the newest fanciest optimization algorithm,...

#### Random search

$$\begin{aligned} & \text{pick a random direction } \mathbf{d}_t \in \mathbb{R}^d \\ & \gamma := \mathop{\mathrm{argmin}}_{\gamma \in \mathbb{R}} f(\mathbf{x}_t + \gamma \mathbf{d}_t) & \text{(line-search)} \\ & \mathbf{x}_{t+1} := \mathbf{x}_t + \gamma \mathbf{d}_t \end{aligned}$$

## **Convergence Rate for Derivative-free Random Search**

**Theorem:** Converges same as gradient descent - up to a slow-down factor d.

**Proof.** Assume that f is a L-smooth convex, differentiable function. For any  $\gamma$ , by smoothness, we have:

$$f(\mathbf{x}_t + \gamma \mathbf{d}_t) \le f(\mathbf{x}_t) + \gamma \mathbf{d}_t^{\top} \nabla f(\mathbf{x}_t) + \frac{\gamma^2 L}{2} \|\mathbf{d}_t\|^2$$

Minimizing the upper bound, there is a step size  $\bar{\gamma}$  for which

$$f(\mathbf{x}_t + \bar{\gamma}\mathbf{d}_t) \le f(\mathbf{x}_t) - \frac{1}{L} \left( \frac{\mathbf{d}_t^{\top}}{\|\mathbf{d}_t\|} \nabla f(\mathbf{x}_t) \right)^2$$

The step size  $\gamma$  we actually took (based on f directly) can only be better:

$$f(\mathbf{x}_t + \gamma \mathbf{d}_t) \le f(\mathbf{x}_t + \bar{\gamma} \mathbf{d}_t)$$
.

Taking expectations, and using the Lemma  $\mathbb{E}_{\mathbf{r}}(\mathbf{r}^{\top}\mathbf{g})^{2} = \frac{1}{d} \|\mathbf{g}\|^{2}$  for  $\mathbf{r} \sim \text{sphere} \subseteq \mathbb{R}^{d}$ :  $\mathbb{E}[f(\mathbf{x}_{t} + \gamma \mathbf{d}_{t})] \leq \mathbb{E}[f(\mathbf{x}_{t})] - \frac{1}{Ld} \mathbb{E}[\|\nabla f(\mathbf{x}_{t})\|^{2}] .$ 

# Convergence Rate for Derivative-free Random Search

Same as what we obtained for gradient descent, now with an extra factor of d. d can be huge!!!

Can do the same for different function classes, as before

- ightharpoonup For convex functions, we get a rate of  $\mathcal{O}(dL/\varepsilon)$  .
- lacktriangle For strongly convex, we get  $\mathcal{O}(dL/\mu\log(1/arepsilon))$  .

Always d times the complexity of gradient descent on the function class.

credits to Moritz Hardt

#### Without the Linear Search: Two Function Evaluations

Without the line search, when a function can be evaluated at two points per iteration, a gradient estimate can be obtained by

$$g_{\alpha}(\mathbf{x}) = \frac{f(\mathbf{x} + \alpha \mathbf{u}) - f(\mathbf{x})}{2\alpha} \mathbf{u}$$
 (1)

where  $\mathbf{u}$  is a random direction sample from the normal distribution  $N(0, I_d)$ . This can be then used as a gradient estimate in the gradient descent update.

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta_t g_\alpha(\mathbf{x}_t) \tag{2}$$

A simplified analysis: We show a preliminary analysis of the algorithm at the limit of  $\alpha \to 0$ . Assuming the function f is differentiable, we have,

$$g_0(\mathbf{x}) = \left(\nabla f(\mathbf{x})^\top \mathbf{u}\right) \mathbf{u} \tag{3}$$

# Simplified Analysis of the Two Function Evaluation

First note that the gradient estimate is unbiased, i.e.,  $\mathbb{E}[g_0(\mathbf{x})] = \nabla f(\mathbf{x})$ . Hence, this Eq. (2) is equivalent to stochastic gradient descent. Next, note second moment can be bounded as follows:

$$\mathbb{E}[\|g_0(\mathbf{x})\|^2] = \mathbb{E}[\|\mathbf{u}\|^2 \nabla f(\mathbf{x})^\top \mathbf{u} \mathbf{u}^\top \nabla f(\mathbf{x})] = \nabla f(\mathbf{x})^\top \mathbb{E}[\|\mathbf{u}\|^2 \mathbf{u} \mathbf{u}^\top] \nabla f(\mathbf{x}).$$

For normal distribution  $N(0, I_d)$ , we have  $\mathbb{E}[\|\mathbf{u}\|^2\mathbf{u}\mathbf{u}^{\top}] \leq (d+3)I_d$ , hence we can write the second moment as:

$$\mathbb{E}[\|g_0(\mathbf{x})\|^2] \leqslant (d+3)\|\nabla f(\mathbf{x})\|^2.$$

In the case of stochastic gradient descent, we have studied the convergence with stochastic gradient oracle  $E[\|g_t\|^2] \leq B$ . Here, we have a *better* control on the second moment which we use in the following analysis.

#### **Analysis of the Two Function Evaluation Method**

For a L-smooth convex function f with optimum at  $\mathbf{x}_*$ , using the vanilla analysis,

$$\mathbb{E}[\|\mathbf{x}_{t+1} - \mathbf{x}_*\|^2 = \|\mathbf{x}_t - \mathbf{x}_*\|^2 - 2\eta_t \mathbb{E}[g_0(\mathbf{x}_t)]^\top (\mathbf{x}_t - \mathbf{x}_*) + \eta_t^2 \mathbb{E}[\|g_0(\mathbf{x}_t)\|^2]$$

$$\leq \|\mathbf{x}_t - \mathbf{x}_*\|^2 - 2\eta_t \mathbb{E}[\nabla f(\mathbf{x}_t)]^\top (\mathbf{x}_t - \mathbf{x}_*) + (d+3)\eta_t^2 \|\nabla f(\mathbf{x}_t)\|^2.$$

With smoothness, we have  $\|\nabla f(\mathbf{x}_t)\|^2 \leq 2L(f(\mathbf{x}_t) - f(\mathbf{x}_*))$ . With convexity, we have  $f(\mathbf{x}_t) - f(\mathbf{x}_*) \leq \nabla f(\mathbf{x}_t)^\top (\mathbf{x}_t - \mathbf{x}_*)$ .

$$\mathbb{E}[\|\mathbf{x}_{t+1} - \mathbf{x}_*\|^2] \le \|\mathbf{x}_t - \mathbf{x}_*\|^2 - 2\eta_t(1 - (d+3)\eta_t L)(f(\mathbf{x}_t) - f(\mathbf{x}_*))$$

$$2\eta_t(1 - (d+3)\eta_t L)(f(\mathbf{x}_t) - f(\mathbf{x}_*)) \le \|\mathbf{x}_t - \mathbf{x}_*\|^2 - \mathbb{E}[\|\mathbf{x}_{t+1} - \mathbf{x}_*\|^2],$$

$$\sum_{t=0}^{T} 2\eta_{t} (1 - (d+3)\eta_{t}L) \mathbb{E}[(f(\mathbf{x}_{t}) - f(\mathbf{x}_{*}))] \leq ||\mathbf{x}_{0} - \mathbf{x}_{*}||^{2} - \mathbb{E}[||\mathbf{x}_{t+1} - \mathbf{x}_{*}||^{2}].$$

With a stepsize  $\eta_t = \eta < \frac{1}{(d+3)L}$ , we get a rate of convergence of  $\mathcal{O}(d/T)$ .

# (Extra) Beyond Simplified Analysis

For the general case of gradient estimate  $g_{\alpha}(\mathbf{x})$ , a similar analysis can be following by computing the expectation and second moment as

$$\mathbb{E}[g_{\alpha}(\mathbf{x})] = \nabla f(\mathbf{x}) + \mathcal{O}(\alpha),$$

$$\mathbb{E}[\|g_{\alpha}(\mathbf{x})\|^{2}] \leq \mathcal{O}(\|\nabla f(\mathbf{x})\|^{2} + \alpha^{2}).$$

The analysis can be followed as before with the correction terms involving  $\alpha$ . For a careful analysis, refer to the paper [DJWW15].

## **Applications for Derivative-free Random Search**

#### **Applications**

- can be used for Reinforcement learning
- memory and communication advantages: never need to store a gradient
- hyperparameter optimization, and other difficult problems like discrete optimization problems
- finding adversarial examples

## **Reinforcement Learning**

$$\mathbf{s}_{t+1} = f(\mathbf{s}_t, \mathbf{a}_t, \mathbf{e}_t)$$
.

where  $s_t$  is the state of the system,  $a_t$  is the control action, and  $e_t$  is some random noise. We assume that f is fixed, but unknown.

We search for a control 'policy'

$$\mathbf{a}_t := \pi(\mathbf{a}_1, \dots, \mathbf{a}_{t-1}, \mathbf{s}_0, \dots, \mathbf{s}_t)$$
.

which takes a trajectory of the dynamical system and outputs a new control action. Want to maximize overall reward

$$\max_{\mathbf{a}_t} \mathbb{E}_{\mathbf{e}_t} \Big[ \sum_{t=0}^{N} R_t(\mathbf{s}_t, \mathbf{a}_t) \Big]$$
s.t.  $\mathbf{s}_{t+1} = f(\mathbf{s}_t, \mathbf{a}_t, \mathbf{e}_t)$ 
(so given)

Examples: Simulations, Games (e.g. Atari), Alpha Go

# **Chapter XI.2**

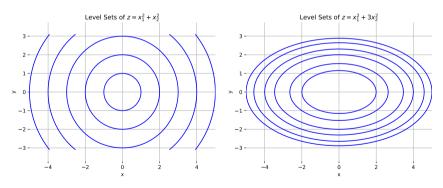
# **Adaptive Methods**

## Some problems with GD

Conside the following function:

$$f(\mathbf{x}) = \begin{bmatrix} \mathbf{x}_1 & \mathbf{x}_2 \end{bmatrix} \begin{bmatrix} \mu & 0 \\ 0 & L \end{bmatrix} \begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \end{bmatrix}$$

The level sets of the functions for different choices of  $\mu$  and L can be seen as follows:



(a) Circular level sets when  $L, \mu = 1$ 

(b) Elliptical level sets when  $L, \mu = 3, 1$ 

#### Some problems with GD

- For skewed functions,  $\mu << L$ , the level sets are ellipses, in these cases GD can be very slow along some directions.
- ightharpoonup A solution to this problem is to use a preconditioner P.

#### **Preconditioned Gradient Descent**

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta_t P_t \nabla f(\mathbf{x}_t),$$

for some semi-definite matrix  $P_t$  which is often called as preconditioner.

Intuition: Consider a function  $g(\mathbf{y}) = f(R\mathbf{y})$ , for some matrix R, now gradient descent on g is given by:

$$\mathbf{y}_{t+1} = \mathbf{y}_t - \eta_t \nabla g(\mathbf{y}_t) = \mathbf{y}_t - \eta_t R^\top \nabla f(R\mathbf{y}_t),$$

$$R\mathbf{y}_{t+1} = R\mathbf{y}_t - \eta_t R R^\top \nabla f(R\mathbf{y}_t)$$

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta_t R R^\top \nabla f(\mathbf{x}_t).$$

Hence, we can choose  $P_t = RR^{\top}$ , and the preconditioned gradient descent is equivalent to the gradient descent on the function  $g(\mathbf{y}) = f(R\mathbf{y})$ .

## How to choose the preconditioner?

- ▶ The optimal preconditioner is the inverse of the Hessian,  $P = (\nabla^2 f(\mathbf{x}))^{-1}$  (similar to the Newton method).
- However, for a function f in d-dimensional space, the Hessian is a  $d \times d$  matrix, making its inversion computationally expensive and often impractical.
- As a result, practical approaches rely on approximations, such as using a diagonal matrix to estimate the Hessian.

#### **Adagrad**

Adagrad is an adaptive variant of SGD.

Pick a stochastic gradient 
$$\mathbf{g}_t$$
. For all  $i$ , 
$$\text{Update } [G_t]_i := \sum_{s=0}^t ([\mathbf{g}_s]_i)^2$$
 
$$\text{Diagonal preconditioner: } P_t = \begin{bmatrix} \sqrt{[G_t]_1} & 0 & \cdots & 0 \\ 0 & \sqrt{[G_t]_2} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \sqrt{[G_t]_d} \end{bmatrix}^{-1},$$
 
$$\mathbf{x}_{t+1} := \mathbf{x}_t - \gamma \, P_t \, \mathbf{g}_t, \implies [\mathbf{x}_{t+1}]_i := [\mathbf{x}_t]_i - \frac{\gamma}{\sqrt{[G_t]_i}} [\mathbf{g}_t]_i.$$

EPFL Optimization for Machine Learning CS-439

## **Adagrad**

- chooses an adaptive, coordinate-wise learning rate
- ▶ To select the preconditioner, intutively first the Hessian is approximated by the sum of gradient outer product matrices,  $G_t = \sum_{s=0}^t \mathbf{g}_s \mathbf{g}_s^{\mathsf{T}}$  and then only the diagonal enteries are used.
- strong performance in practice; theoretical guarantees for convergence in convex setting with better performance than GD especially for sparse gradients [DHS11]
- ▶ However, Adagrad often leads to diminishing learning rates as the aggregates square gradients,  $G_t$ , quickly become large.
- Variants: Adadelta, RMSprop, Adam

## **RMSprop**

RMSprop is a variant of Adagrad that uses an exponential moving average with a parameter  $\beta$  of the squared gradients instead of the sum.

Pick a stochastic gradient  $\mathbf{g}_t$ . For all i,  $\mathsf{Update}\ [G_t]_i := (\beta)[G_{t-1}]_i + (1-\beta)([\mathbf{g}_t]_i)^2$   $\mathbf{x}_{t+1} := [\mathbf{x}_t]_i - \frac{\gamma}{\sqrt{|G_t|_i}}[\mathbf{g}_t]_i.$ 

#### Momentum SGD

Before presenting Adam, lets discuss another momentum variant of SGD a classical technique propsed by Polyak in 1964 to accelerate gradient descent.

```
pick a stochastic gradient \mathbf{g}_t \mathbf{m}_{t+1} := \beta \mathbf{m}_t + (1-\beta)\mathbf{g}_t \qquad \qquad \text{(momentum term)} \mathbf{x}_{t+1} := \mathbf{x}_t - \gamma \mathbf{m}_{t+1}
```

(standard choice of  $\mathbf{g}_t := \nabla f_j(\mathbf{x}_t)$  for sum-structured objective functions  $f = \sum_j f_j$ )

- momentum from previous gradients
- ▶ is a variant of the Nesterov acceleration seen before
- key element of deep learning optimizers, necessary for top accuracy

#### **Polyak Momentum and Nesterov Acceleration**

#### Polyak Momentum

At iteration t,

$$\mathbf{m}_{t+1} = \frac{\beta \mathbf{m}_t}{\eta} + \eta \nabla f(\mathbf{x}_t),$$
  
 $\mathbf{x}_{t+1} = \mathbf{x}_t - \mathbf{m}_{t+1}$ 

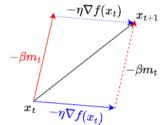


Figure: Aggregating the momentum and the gradient

#### **Nesterov Acceleration**

At iteration t,

$$\mathbf{m}_{t+1} = \frac{\beta \mathbf{m}_t}{\beta \mathbf{m}_t} + \eta \nabla f(\mathbf{x}_t - \beta \mathbf{m}_t),$$
  
$$\mathbf{x}_{t+1} = \mathbf{x}_t - \mathbf{m}_{t+1}$$

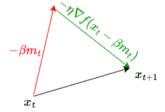


Figure: 'Look ahead' gradient. Add the momentum and take a gradient step from the new position.

# **Convergence of Momentum and Nesterov Acceleration**

- For quadratic functions, both methods converge at an accelerated rate of  $\exp\{-t/\sqrt{\kappa}\}$ , where  $\kappa$  is the condition number of the function.
- For strongly convex functions, Nesterov acceleration converges at a rate of  $\exp\{-t/\sqrt{\kappa}\}$ , while Polyak momentum does not provably converge at a faster rate.

#### Adam

#### Adam is a momentum variant of Adagrad

```
pick a stochastic gradient \mathbf{g}_t \mathbf{m}_t := \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t \qquad \qquad \text{(momentum term)} [\mathbf{v}_t]_i := \beta_2 [\mathbf{v}_{t-1}]_i + (1 - \beta_2) ([\mathbf{g}_t]_i)^2 \quad \forall i \quad \text{(2nd-order statistics)} [\mathbf{x}_{t+1}]_i := [\mathbf{x}_t]_i - \frac{\gamma}{\sqrt{[\mathbf{v}_t]_i}} [\mathbf{m}_t]_i \quad \forall i
```

- faster forgetting of older weights
- momentum from previous gradients (see acceleration)
- ightharpoonup (simplified version, without correction for initialization of  $\mathbf{m}_0, \mathbf{v}_0$ )
- strong performance in practice, e.g. for self-attention based networks like transformers

- ▶ While RMSprop and Adam solve the problem of diminishing learning rates, they can lead to convergence issues particularly due to the "short-term" memory of the exponential moving average of the squared gradients.
- ► EMA of the squared gradients approximately limits the update to only a few past gradients. Hence, particularly in the case where a few minibatches provide large gradients, their influence dies out quickly due to the EMA.

Problem Setup: Consider the functions  $f_t(x) = \begin{cases} Cx & \text{for } t \mod 3 = 1, \\ -x & \text{otherwise.} \end{cases}$  , for C > 2

over the domain  $\mathcal{F}=[-1,1]$ . We are interested in the following quantity (regret) over the course of optimization

$$R_T = \frac{1}{T} \sum_{t=1}^{T} (f_t(x_t)) - \frac{1}{T} \min_{x \in \mathcal{F}} \sum_{i=1}^{T} f_t(x)$$

Note that x=-1 gives the minimum regret, i.e.,  $\min_{x\in\mathcal{F}}\sum_{i=1}^T f_t(x)$ . However, Adam provably converges to highly suboptimal x=+1 [RKK18].

Note the gradient of the function  $f_t(x)$  is given by:

$$\nabla f_t(x) = \begin{cases} C \text{ for } t \mod 3 = 1, \\ -1 \text{ otherwise.} \end{cases},$$

i.e., one large gradient every 3 iterations. However, the EMA of squared gradients nullify this large gradient. It can be shown that for the update of Adam,

#### Lemma

Consider the Adam algorithm with appropriate choice of  $\beta_1$  and  $\beta_2$  starting with  $x_1 = 1$ , then

$$x_t = 1$$
 for every  $t$  such that  $t \mod 3 = 1$ 

Proof can be done by induction following [RKK18].

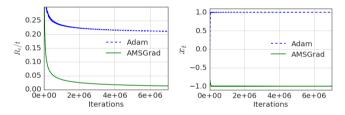


Figure: Non-convergence of ADAM on a simple one dimensional convex problem  $f_t(x)$ . The second plot shows that ADAM converges to x=1. AMSGrad is the algorithm proposed in [RKK18]. Image is taken from [RKK18].

# Chapter XI.3

#### **Efficient variants of SGD**

# SignSGD

Only use the sign (one bit) of each gradient entry: SignSGD is a communication efficient variant of SGD.

pick a stochastic gradient 
$$\mathbf{g}_t$$
 
$$[\mathbf{x}_{t+1}]_i := [\mathbf{x}_t]_i - \gamma_t \, sign([\mathbf{g}_t]_i) \qquad \forall i$$

(with possible rescaling of  $\gamma_t$  with  $\|\mathbf{g}_t\|_1$ )

- communication efficient for distributed training
- convergence issues

## **ClippedSGD**

Clip the gradients to a predefined maximum length c>0.

$$\begin{aligned} & \text{pick a stochastic gradient } \mathbf{g}_t \\ & \mathbf{x}_{t+1} := \mathbf{x}_t - \gamma_t \, clip_c(\mathbf{g}_t) \end{aligned} \qquad \forall i$$

with 
$$clip_c(\mathbf{g}) := \min(1, \frac{c}{\|\mathbf{g}\|}) \cdot \mathbf{g}$$

- used to avoid instabilities in deep learning training (such as LLMs)
- used with differential privacy (adding noise of comparable magnitude)
- convergence non-trivial

# **Bibliography**

John Duchi, Elad Hazan, and Yoram Singer.

Adaptive subgradient methods for online learning and stochastic optimization.

Journal of machine learning research, 12(7), 2011.

John C Duchi, Michael I Jordan, Martin J Wainwright, and Andre Wibisono. Optimal rates for zero-order convex optimization: The power of two function evaluations.

IEEE Transactions on Information Theory, 61(5):2788–2806, 2015.

Sashank J Reddi, Satyen Kale, and Sanjiv Kumar. On the convergence of adam and beyond. In *ICLR*, 2018.